

# CHAPTER 8

## Programming in Assembly Part 4: I/O Programming

# Relative Speed

- I/O devices are sometimes mechanical devices (e.g., solenoids, relays, etc.) that take a long time to perform an action.
- The computer performs operations orders of magnitude faster than the I/O devices.
- **Synchronization**: The CPU must wait for the I/O device to finish each command before issuing the next.

# Asynchronous Events

- The events that determine when an input device has data available or when an output device needs data are **independent** of the CPU.
- I/O programming, therefore, requires "**hand-shaking**" between the CPU and the I/O device to coordinate the transfer so that data is transferred reliably.

# Time Behavior

<i>Data Rate</i>	<i>Pattern</i>	<i>Increasing Time →</i>
<b>Low</b>	Random	
	Periodic	
<b>High</b>	Random	
	Periodic	

# I/O Data Ports

**Output Device:** Wait for status port to indicate that device is ready to accept data, then copy data to output data port.

- Example: UART Transmit Holding Register

**Input Device:** Wait for status port to indicate that data is available, then copy data from the input data port.

- Example: UART Receive Buffer Register

# I/O Status Ports

(Example: UART Line Status Register)

- **Output Device:** Status port contains a single bit that indicates whether or not the device is ready to accept new data from the data port.

7	6	5	4	3	2	1	0
Receiver Error	Trans. Empty	THRE	Break	Framing Error	Parity Error	Overrun Error	Data Ready

THRE=1 when device is ready to receive new data; writing new data into the output data port resets this bit to 0.

- **Input Device:** Status port contains a single bit that indicates when new data is available to be read from the data port.

7	6	5	4	3	2	1	0
Receiver Error	Trans. Empty	THRE	Break	Framing Error	Parity Error	Overrun Error	Data Ready

Data Ready will go to 1 when new data is available in the input data port; reading that data from the data port resets this bit to 0.

# I/O Control Ports

(Example: UART Line Control Register)

- Used to control operational features of the device.

7	6	5	4	3	2	1	0
DLAB	Set Break	Stick Parity	Even Parity Select	Parity Enable	Number of Stop Bits	Word Length Select Bits	

# Accessing I/O Ports

I/O Mapped:  
(Desktop Processors)

Processor has an I/O address bus separate from the memory address bus used by special instructions for transferring data to and from an I/O port.

Memory-Mapped:  
(Embedded Processors)

I/O devices and memory both share parts of the same memory address space. I/O devices are accessed like regular memory locations.



# Three Strategies

- Polled Waiting Loops
- Interrupt-driven I/O
- Direct Memory Access (DMA)

# Polled Waiting Loop

## (Example: Output data to UART)

```
// THRE=1 if transmitter holding register empty
#define THRE    (1 << 5)

#define DATAPORT(base)    *(base + 0)
#define STATUSPORT(base) *(base + 24)

void Output(uint8_t data, uint8_t *uart_base)
{
    while (STATUSPORT(uart_base) & THRE) == 0)
    {
        // do nothing (wait for device to become
ready)
    }
    DATAPORT(uart_base) = data ;
}
```

# Polled Waiting Loop

## (Example: Input data from UART)

```
// DataReady = 1 if data available to be read
#define DATAREADY (1 << 0)
```

```
#define DATAPORT(base)      *(base + 0)
#define STATUSPORT(base)   *(base + 24)
```

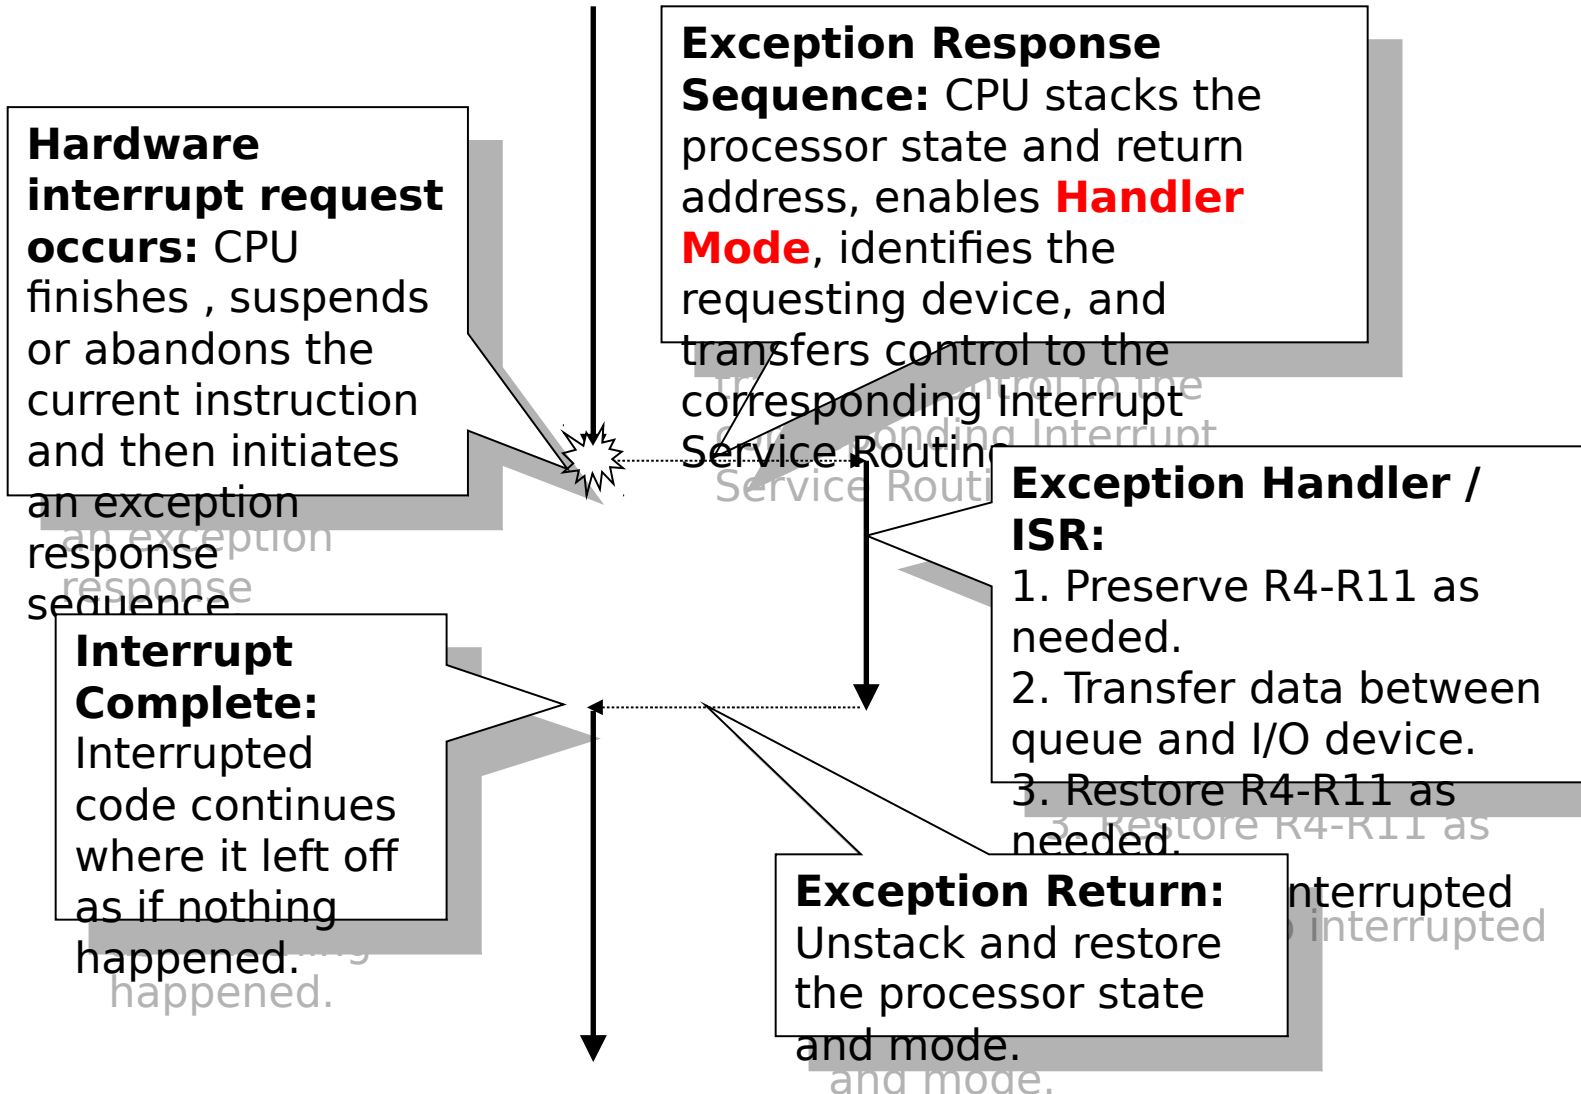
```
uint8_t Input(uint8_t *uart_base)
```

```
{
    while (STATUSPORT(uart_base) & DATAREADY) == 0)
    {
        // do nothing (wait for device to become
ready)
    }
    return DATAPORT(uart_base) ;
}
```

What if urgent data arrives on another input port while executing this

# EXCEPTIONS AND INTERRUPTS

# Interrupt Processing



# Exceptions

## Definitions:

- Any condition that needs to halt the normal sequential flow of instruction execution.
- Reset
- SVC Supervisor Call  
(Software Interrupt)
- Fault  
E.g., undefined opcode
- Interrupts

# Exceptions

Each exception has:

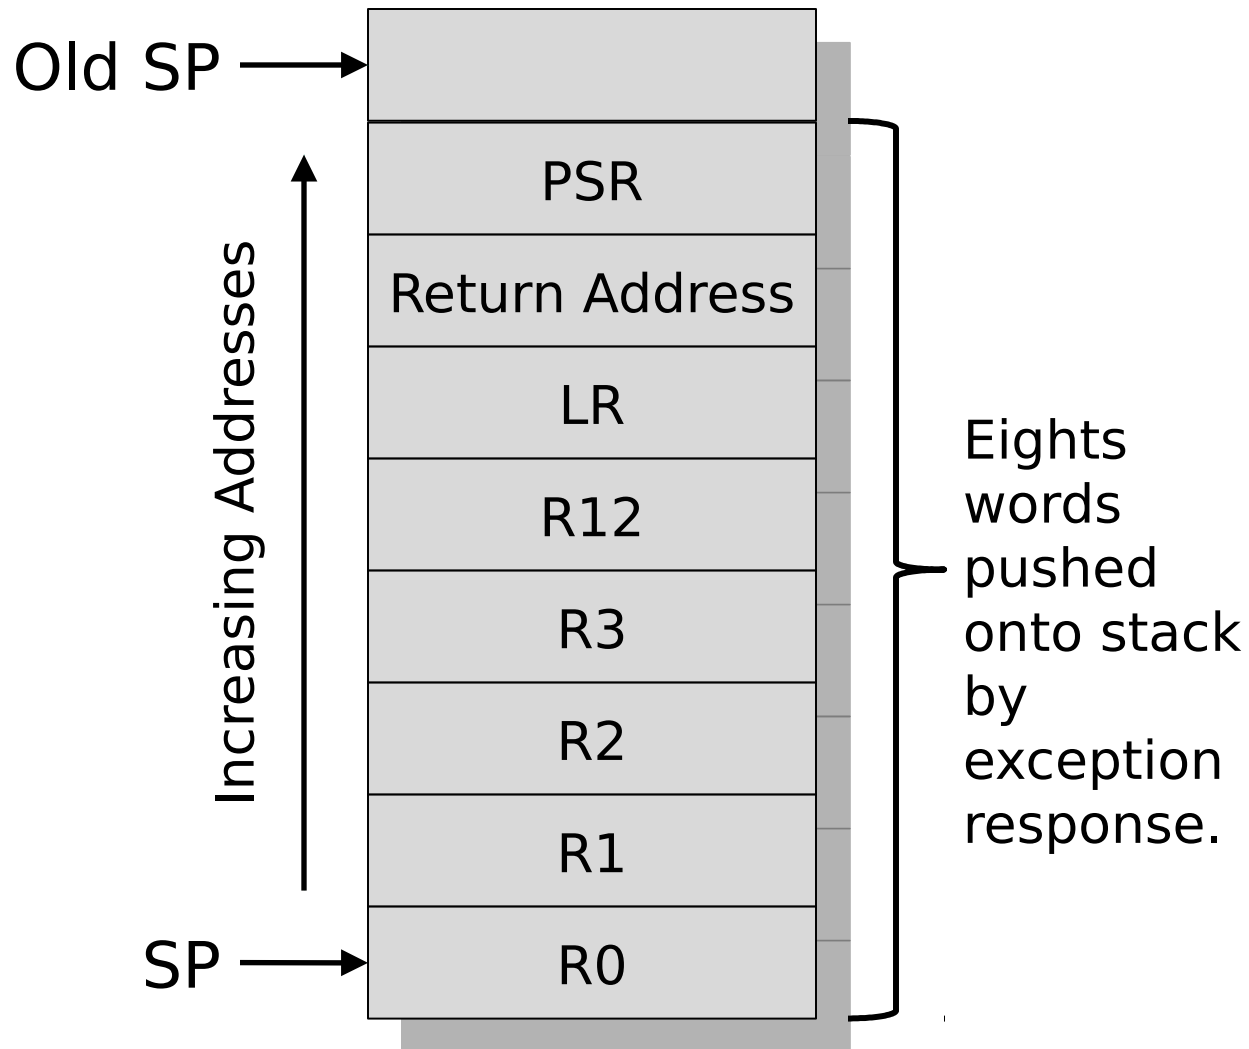
1. An exception number
2. A priority level
3. An exception handler routine
4. An entry in the vector table  
(Address of associated handler routine)

# Exception Response

1. Processor state (8 words) stored on stack:  
xPSR, Return Address, LR, R12, R3 - R0  
*Allows a regular C function to be an ISR!*
2. Processor switched (from **Thread Mode**) to **Handler Mode**  
recorded in “T” Bit (#24) of pushed xPSR.
3. PC  $\equiv$  vector table[ exception # ]



# Interrupt Stacking



# Exception Handlers

- Definition:

An exception handler is a software routine that is executed when a specific exception condition occurs.

Note: Most, but not all, exception handlers return to the previous code.

# Exception Return

Exception return occurs when in *Handler Mode* and one of the following instructions is executed:

1. POP/LDM includes the PC, or
2. LDR with PC as the destination, or
3. BX with any register as the source

“Unstacks” (pops and restores) registers and return address (into PC) from the stack.

- Restores T-bit (in xPSR) and thus previous processor mode.
- Resumes suspended multi-cycle instruction (if any), or restarts abandoned instruction (if any)

# LATENCY

The passage of time from the moment an event occurs until the processor executes code to process the event.

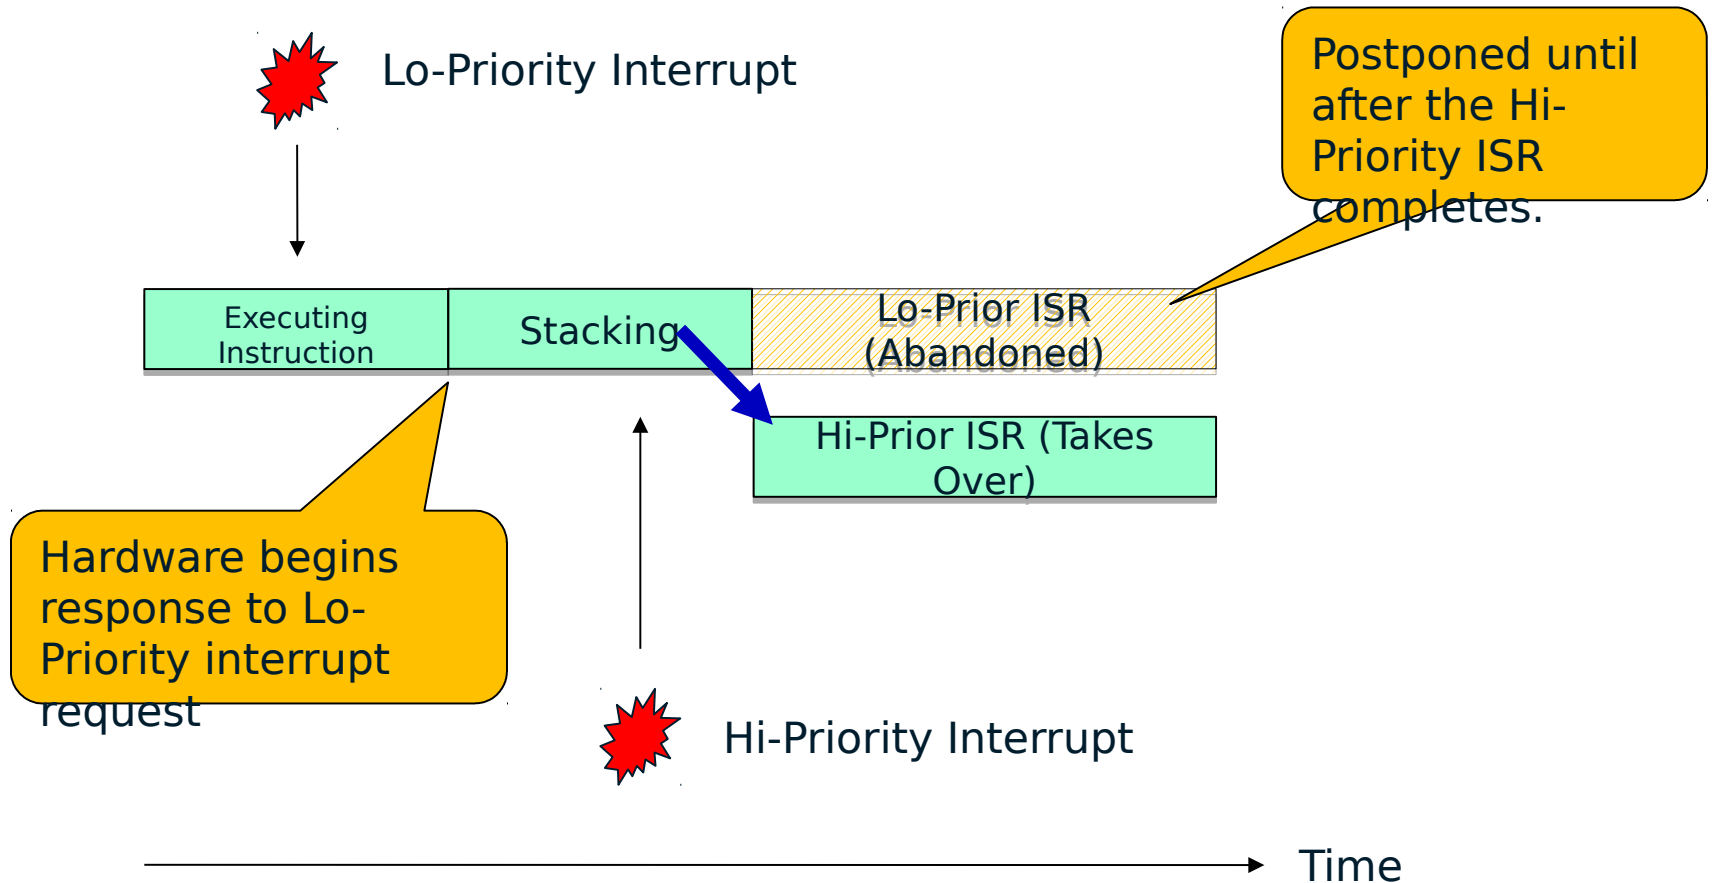
# ARM Latency Reduction

1. **Abandoning or Suspending** the current instruction to start the exception response sooner.
2. **Late Arrival Processing**: Allowing a high-priority event to take over the exception response of a low-priority event.
3. **Tail-Chaining**: Eliminating unnecessary stacking and unstacking.

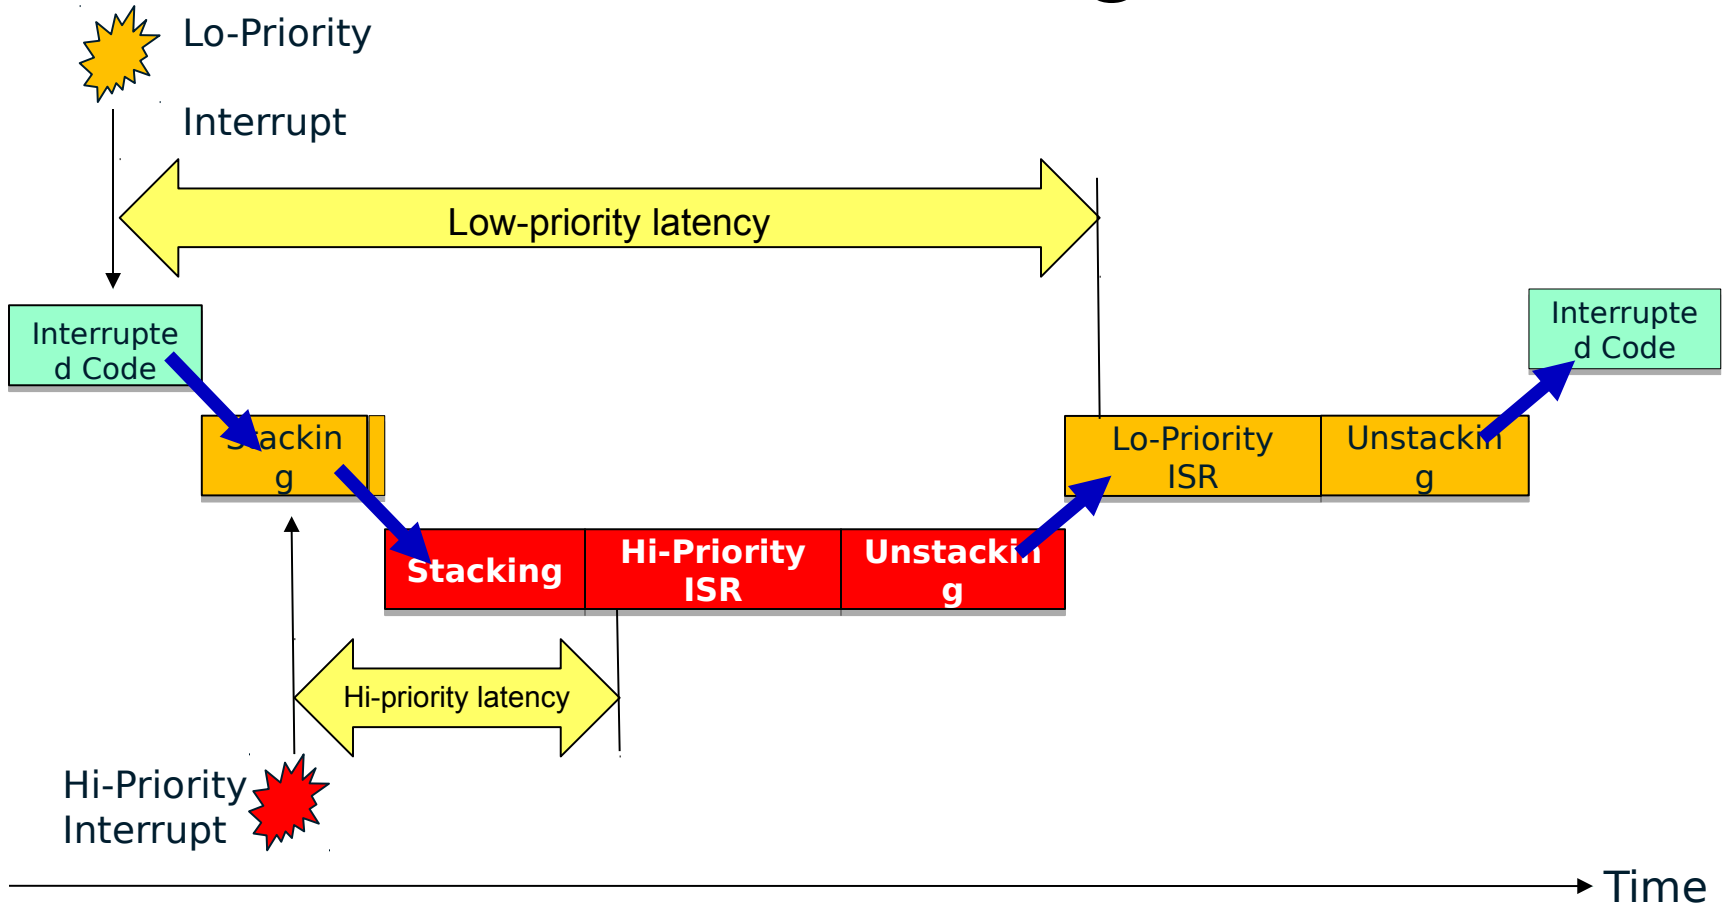
# Suspending/Abandoning Instructions

1. Instructions are allowed to complete if only 1 clock cycle is required
2. Instructions that transfer multiple words to/from memory are suspended  
(LDM, STM, PUSH & POP)
3. All other instructions are abandoned and restarted after the interrupt.

# Late Arrival Processing

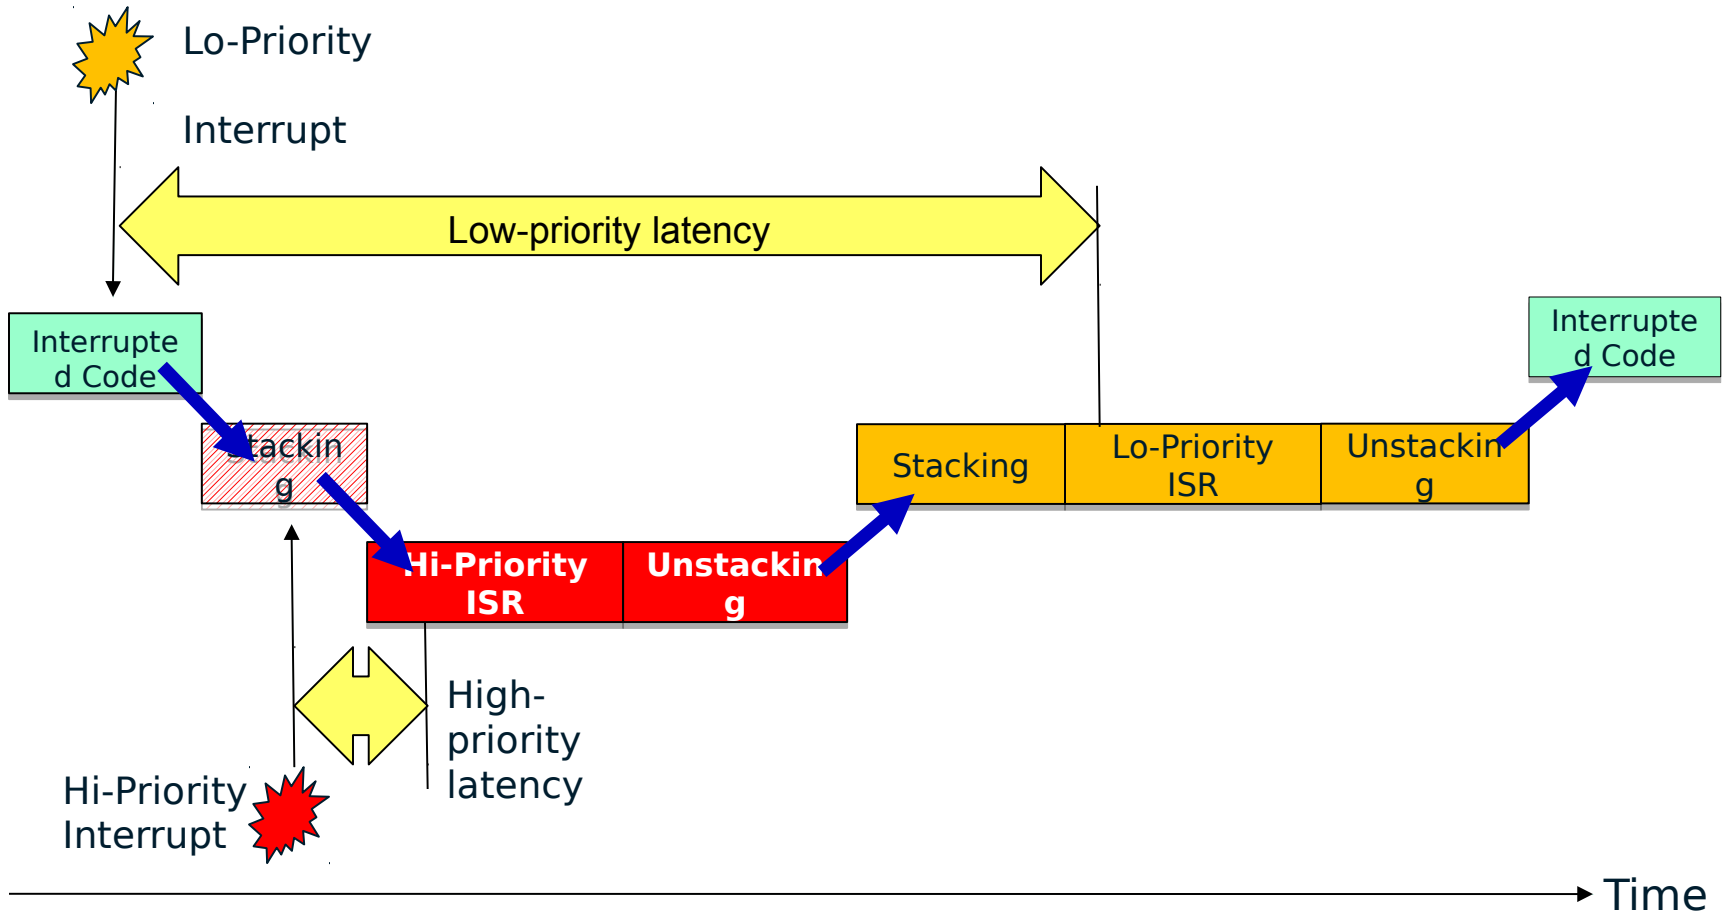


# Without Late Arrival Processing

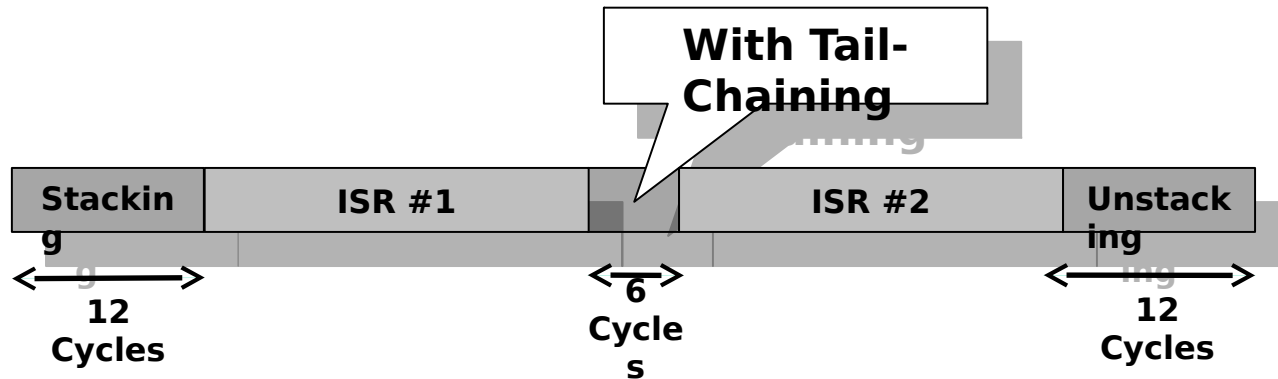
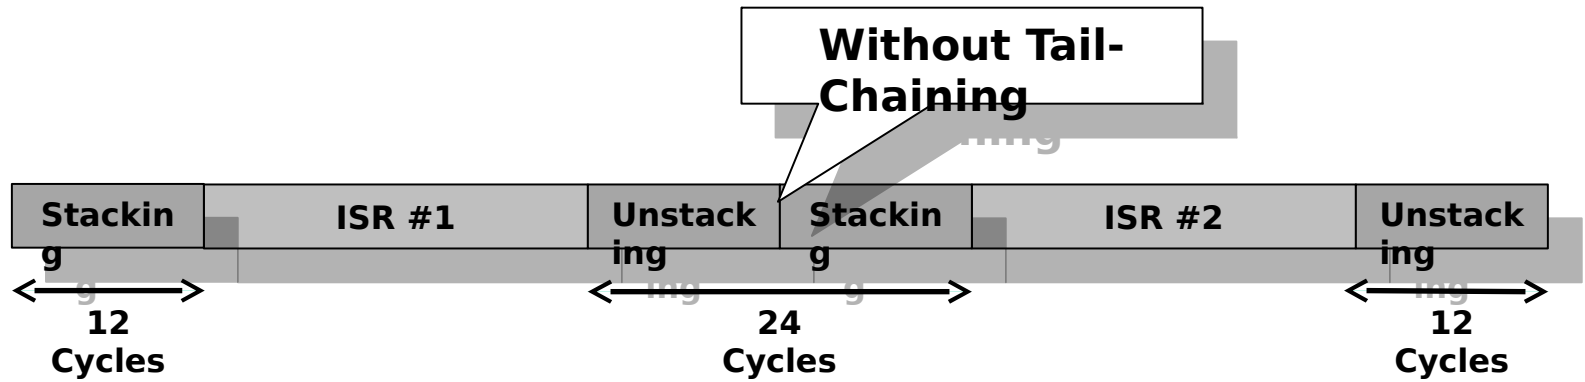




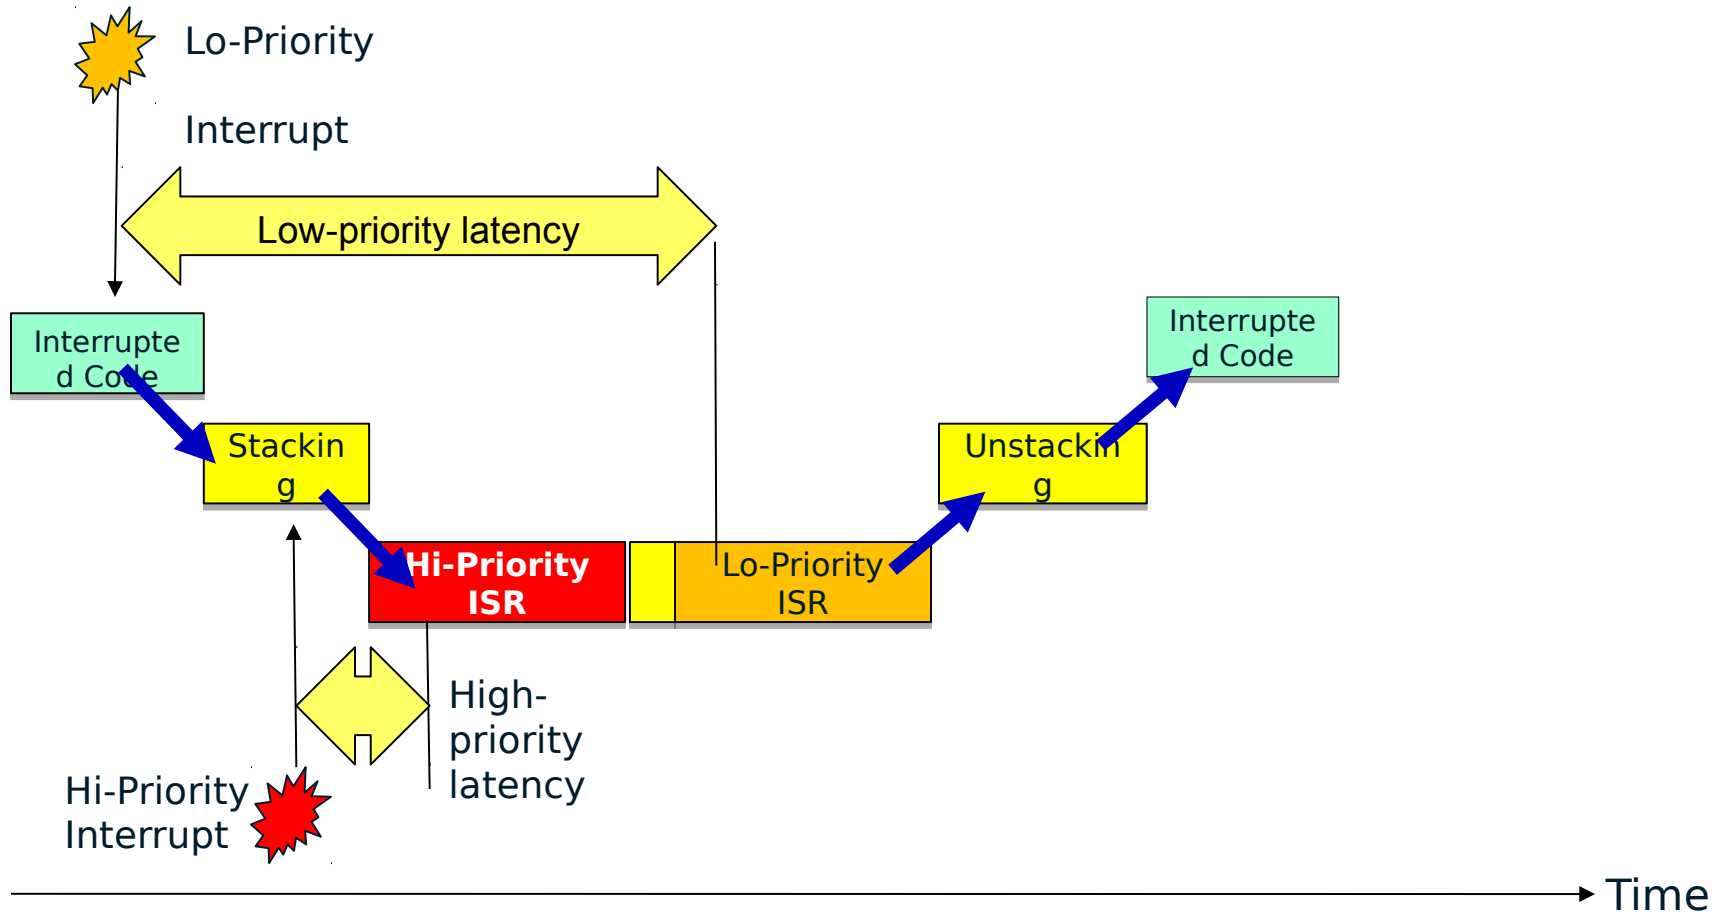
# With Late Arrival Processing



# Tail-Chaining



# With LAP and Tail-Chaining



# EXCEPTION AND INTERRUPT PRIORITIES

# Nested Vectored Interrupt Controller

(Mapped to addresses E000E100-E000ECFF<sub>16</sub>)

*The NVIC provides the ability to:*

1. Individually Enable/Disable interrupts from specific devices.
2. Establishes relative priorities among the various interrupts.

# Exception Priorities and Vector Table Positions

Exception Type	Exception # <sup>1</sup>	IRQ #	Priority	Comment
		n/a		Initial SP value (loaded on reset)
Reset	1		-3 (fixed)	Power up and warm reset
NMI	2		-2 (fixed)	Non-Maskable Interrupt
Hard Fault	3		-1 (fixed)	
Memory Mgmt.	4		Configured through System Handler Priority Registers	
Bus Fault	5			Address/Memory-related faults
Usage Fault	6			Undefined instruction
Reserved	7-10			
SVCall	11			Software Interrupt (SVC instruction)
Debug Monitor	12			
Reserved	13			
PendSV	14			
SysTick	15			System Timer Tick
Ext. Interrupts	16-255	0-239	Configurable	

<sup>1</sup>Vector table entries are ordered by exception number, with the first entry reserved for the initial SP value. IRQ # is the interrupt request number (same as the exception number minus 16).

# Interrupt Priority Registers

- Each interrupt is assigned an 8-bit priority number:
  - Highest Priority: 0
  - Lowest Priority: 255

# Enabling/Disabling Interrupts

CPSIE ; Enable **all** external  
interrupts

CPSID ; Disable **all** external  
interrupts

31

0

NVIC Interrupt Clear Enable Reg.

31

0

NVIC Interrupt Set Enable Reg.

Set a bit to 1 in this  
register to [disable](#)  
interrupts from an  
external I/O device

Set a bit to 1 in this  
register to [enable](#)  
interrupts from an  
external I/O device



# NVIC External Interrupts


0-4	Match Point Timer
5-24	UART0-21
25	Analog Comparator
26-27	Reserved
28	System Fault
29-32	PMU Controller 0-3
33-31	Reserved
14-17	ADC Sequence 0-3



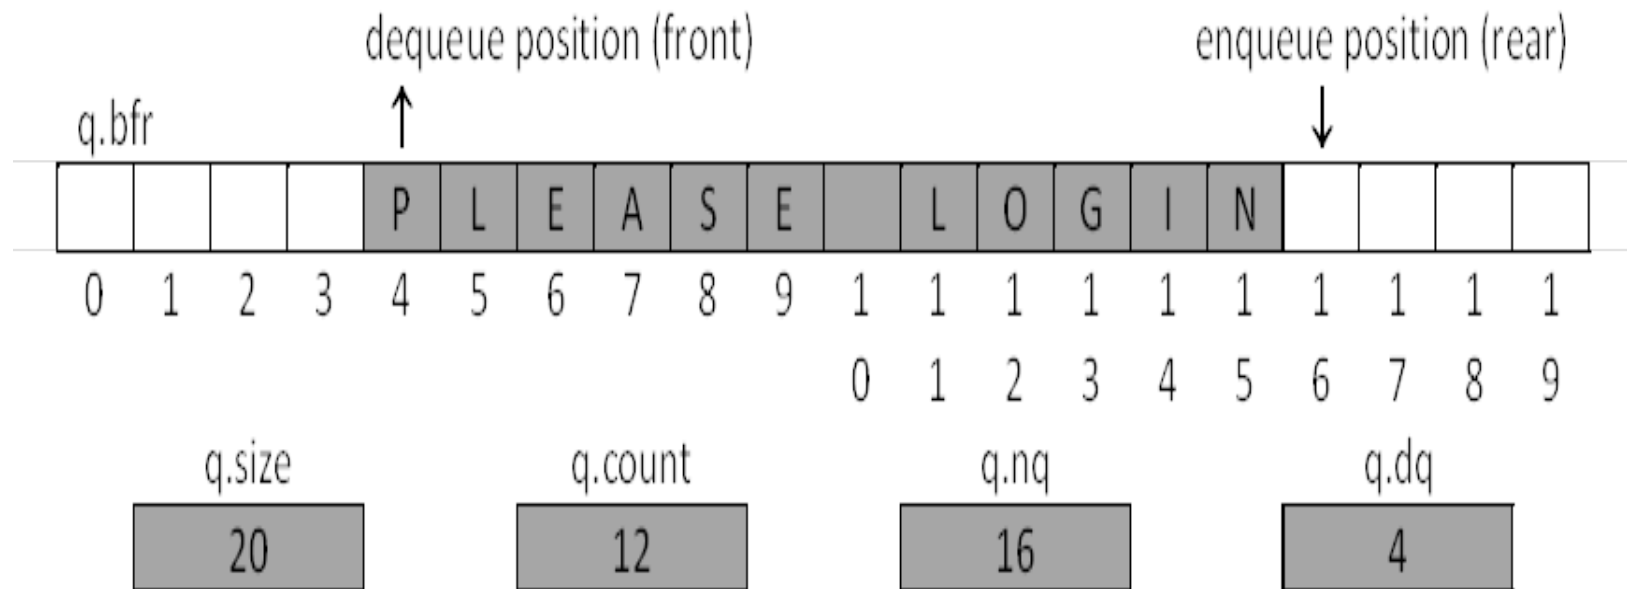
Bit in Interrupt  
Registers

# BUFFERING

# Using a Queue to Decouple Data Arrival from Data Processing



# Interrupt Buffering



# Implementing a Queue in C

```
BOOL Enqueue(Queue *q, uint8_t data)
```

```
{
    BOOL full ;

    disable() ;
    full = q->count == q->size ;
    if (!full)
    {
        q->bfr[q->nq] = data ;
        if (++q->nq == q->size) q->nq = 0 ;
        q->count++ ;
    }
    enable() ;

    return !full ;
}
```

```
BOOL Dequeue(Queue *q, uint8_t *ptr2data)
```

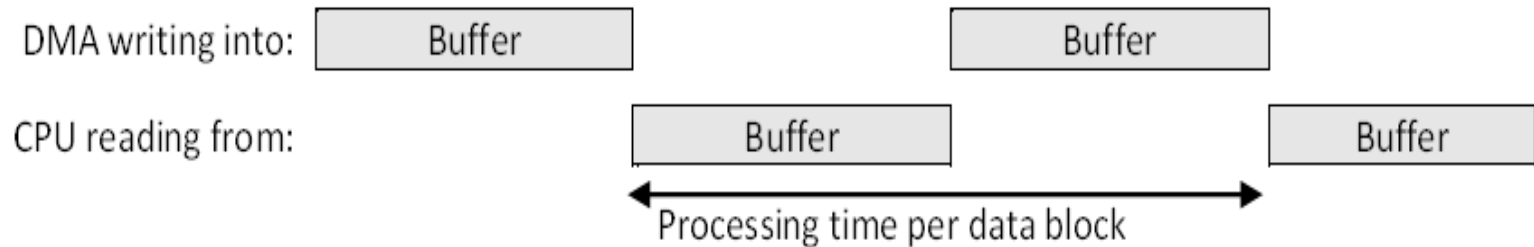
```
{
    BOOL empty ;

    disable() ;
    empty = q->count == 0 ;
    if (!empty)
    {
        *ptr2data = q->bfr[q->dq] ;
        if (++q->dq == q->size) q->dq
            = 0 ;

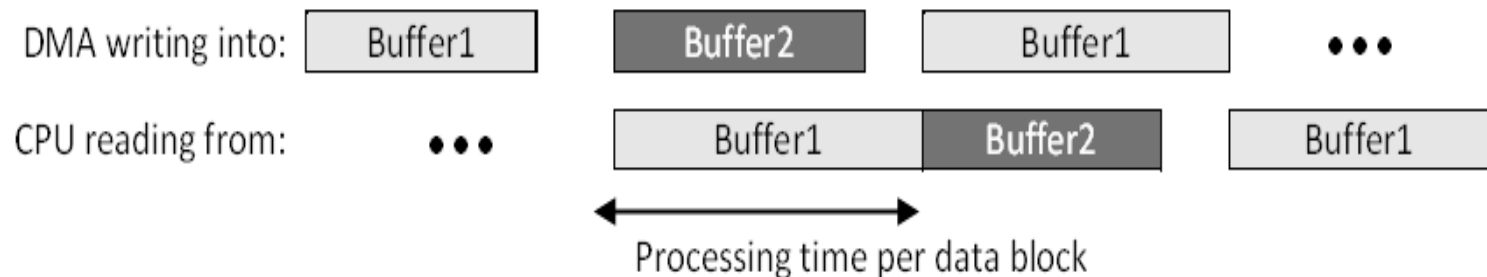
        q->count-- ;
    }
    enable() ;

    return !empty ;
}
```

# Single Buffered DMA



# Double Buffered DMA



# SYNCHRONIZATION, TRANSFER RATE AND LATENCY

# Polled Waiting Loops

- Test device status in a ***waiting loop*** before transferring each data byte.
- **Maximum data rate**: Determined by the time required to execute one iteration of the waiting loop plus the transfer.
- **Latency**: Unpredictable - no guarantee when the program will arrive at the waiting loop.



# Synchronizing data transfers using polled waiting loops.

```
// TXFE (bit 7 of status port) =1 if transmitter holding register empty
```

```
#define TXFE      (1 << 7)
```

```
// RS232 Data Port = rs232 base address + 0
```

```
// RS232 Status Port      = rs232 base address + 24
```

```
void rs232Output(uint8_t *ptr2data, int bytes, volatile uint32_t *rs232base)
{
    for (int byte = 0; byte < bytes; byte++)
    {
        while ((*rs232base + 6) & TXFE) == 0)
        {
            // do nothing (wait for data to arrive)
        }
        *rs232base = (uint32_t) *ptr2data++;
    }
}
```

# Polled waiting loop: Optimized in assembly.

```
EXPORT rs232Output  
rs232Output:
```

```
; R0 contains the parameter 'ptr2data'  
; R1 contains the parameter 'bytes'  
; R2 contains base address of RS232 device
```

```
Repeat:  CBZ      R1,Return    ; more data to send?  
OutWait: LDR      R3,[R2,#0x18] ; get status word  
         TST      R3,#0x80    ; TXFE = 1?  
         BEQ      OutWait     ; loop until ready  
         LDRB     R3,[R0],#1   ; get next byte & update pointer  
         STR      R3,[R2]     ; send data to device  
         SUB      R1,R1,#1    ; decrement byte count  
         B        Repeat      ; do it all again  
Return:  BX       LR          ; return
```

# Cortex-M3 Instruction Clock Cycles<sup>1</sup>

Instructions		Clock Cycles
PUSH, POP, LDM, STM		1 + #regs
SDIV, UDIV		2 - 12
SMLAL, UMLAL		4 - 7
SMULL, UMULL		3 - 5
Unconditional Branch (B, BL, BX)		2 - 4
Conditional Branch	Successful	2 - 4
	Failed	1
LDRD, STRD		3
ADR, MLA, MLS, & all LDR's and STR's		2
All other instructions		1

<sup>1</sup>Assumes memory and processor run at the same speed.

# Data Rate: Polled Waiting Loop

```
EXPORT rs232Output          ; Clock
rs232Output:                ; Cycles

Repeat:  CBZ      R1,Return   ; 1 (Fail)
OutWait:  LDR      R3,[R2,#0x18] ; 2
          TST      R3,#0x80   ; 1
          BEQ      OutWait    ; 1 (Fail)
          LDRB     R3,[R0],#1  ; 2
          STR      R3,[R2]     ; 2
          SUB      R1,R1,#1    ; 1
          B        Repeat     ; 4
```

Clock cycles: 14

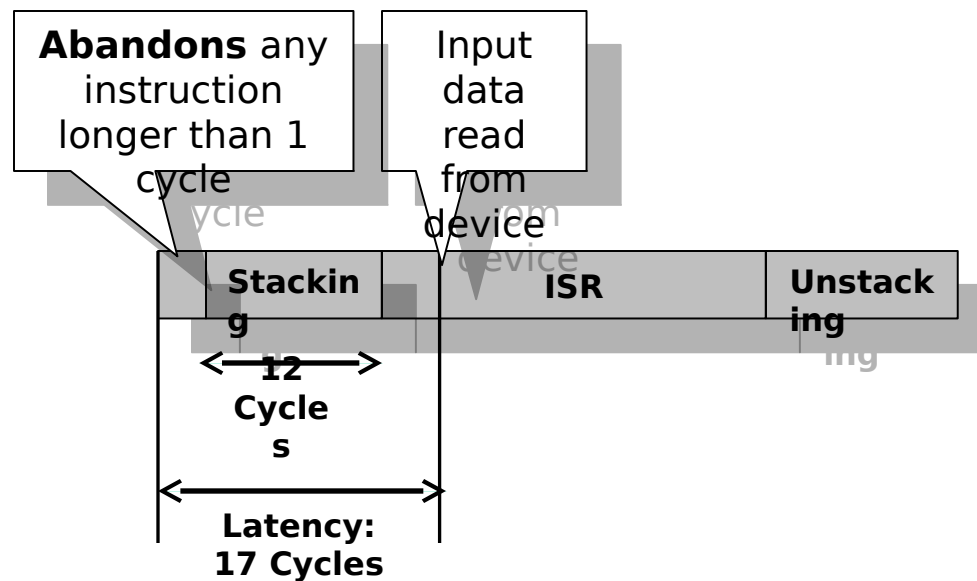
Execution time:  $14 \times 20 \text{ nsec} = 280 \text{ nsec}$

Maximum Data Rate:  $10^9 \div 280 = \mathbf{3.6 \text{ MB/sec}}$

# Data Rate:Interrupt-Driven I/O

Rs232InputHandler:			Clock Cycles
ADR	R0,rs232InpDataPort ; R0 $\equiv$ rs232 device address		2
LDR	R0,[R0] ; R0 $\equiv$ rs232 input data byte		2
LDR	R2,rs232InputNQIndex ; update enqueue index		2
ADD	R2,R2,#1		1
AND	R2,R2,#0x3F ; increment & wrap (0 to 63)		1
STR	R2,rs232InputNQIndex		2
LDR	R3,rs232InputQueue ; R3 $\equiv$ address of buffer		2
STRB	R0,[R3,R2] ; store data into buffer		2
LDR	R2,rs232InputCount ; update #items in queue		2
ADD	R2,R2,#1		1
STR	R2,rs232InputCount		2
BX	LR ; return		4
Tail-chaining time between successive ISRs:			6
Clock cycles:			29
Execution time:			$29 \times 20 \text{ nsec} = 580 \text{ nsec}$
Maximum Data Rate:			$10^9 \div 580 = \mathbf{1.7 \text{ MB/sec}}$

# Interrupt Latency



Total latency: 17 Cycles × 20 nsec = **340 nsec**

# Direct Memory Access

- Requires additional hardware to control data transfers independent of CPU.
- Competes with CPU for control of memory.
- Does not have to wait for current instruction to complete – only the current memory cycle.
  - Latency is thus 1 memory cycle.
  - If memory and I/O are on separate chips: 1 cycle = **60 nsec**
- Data Rate determined by memory speed.
  - 32-bit data bus: 4 bytes / ( $60 \times 10^{-9}$  sec) = **66 MB/Sec**

# Performance Estimates and Relative Cost

	Polled Waiting Loop	Interrupt- Driven I/O	Direct Memory Access
<b>Maximum Transfer Rate</b>	~3.6 MB/sec	~1.7 MB/sec	~66 MB/sec
<b>Best-Case Latency</b>	unpredictable	340 nsec	~60 nsec
<b>Hardware Cost</b>	Least	Low	Moderate
<b>Software Complexity</b>	Low	Moderate	Moderate